



PREMIER

BUSINESS PARTNER


# Docker

## Architectural Impact

*START >*

# Docker Architectural Impact

- Applications
  - Configure
  - Deploy
  - Backup/Restore
- Monitoring
  - (Distributed) Logging
  - Classical Monitoring Systems
- Security
  - Docker Infrastructure
  - Docker Image-Registry
  - Docker Containers



docker

<https://docker.com>

# Application

## Monolith:

- One repository
- One artifact represents application
- One build, one deployment
- One configuration per deployment
- Touches ConfigSource
- Uses Secrets, and knows them
- Centralized logging
- Application in one runtime
- Fix/Deploy whole application
- ...

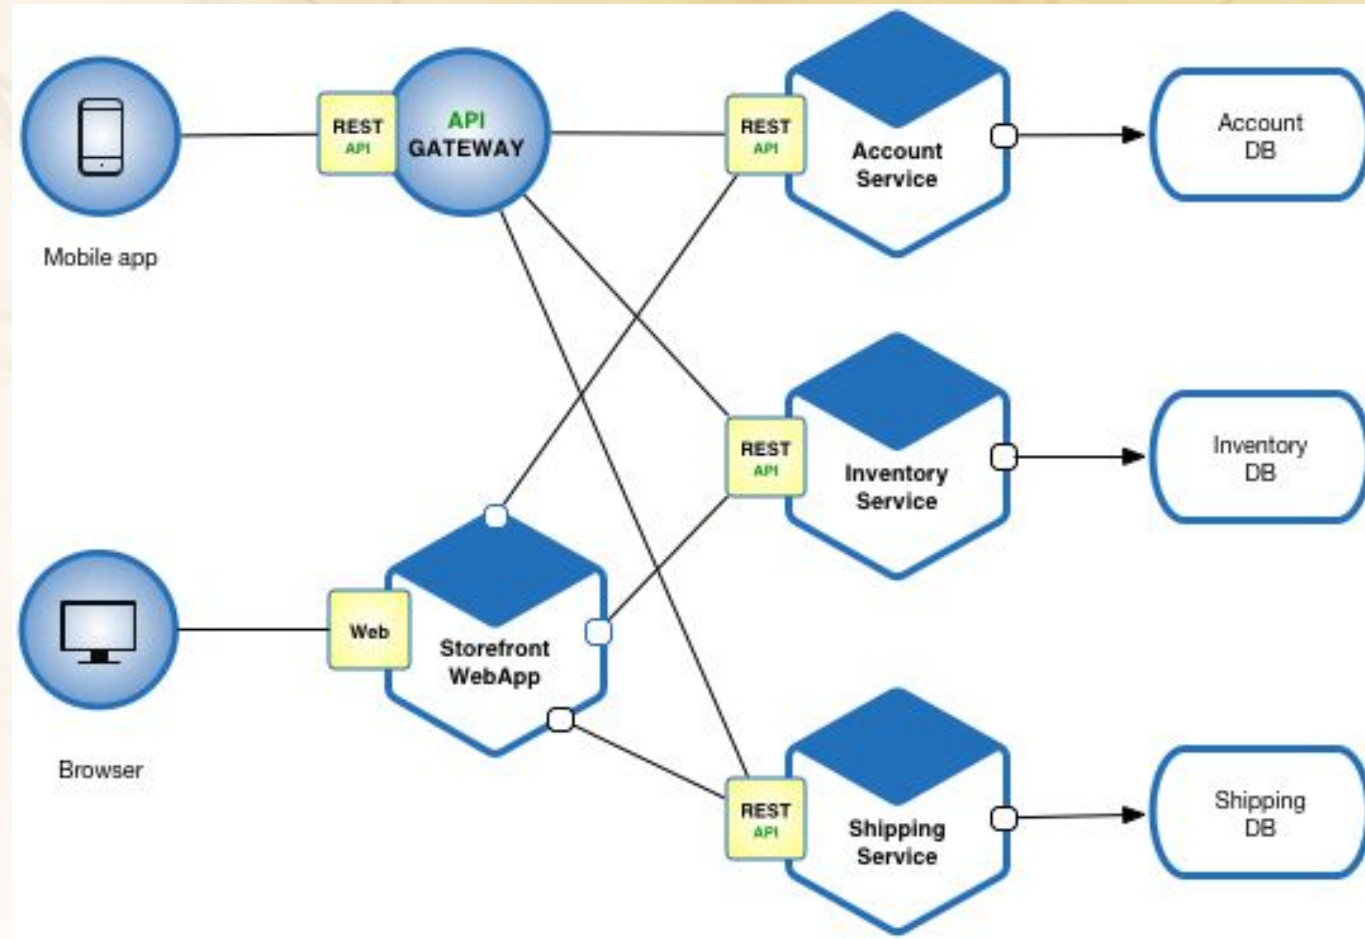
## Microservice:

- N repositories
- N artifacts represents application
- N builds and deployments
- One configuration per N deployments
- No reference to ConfigSource (*MicroProfile Config*)
- Uses Secrets, but doesn't know them
- Distributed Logging (*MicroProfile OpenTracing*)
- Application in N runtimes
- Fix/Deploy single Service
- ...



# Application

- More flexible and complex  
*(but not complicated)*
- Easier to maintain
- Harder to manage
- Application is distributed
- New problems
  - Failures are distributed
  - Logs are distributed
  - Possibility of Cycles
  - Service resilience very important



[https://www.youtube.com/watch?v=TvnZTi\\_gaNc](https://www.youtube.com/watch?v=TvnZTi_gaNc)

# Application - Configure

- Configurations can be provided via:
  - **Environment variables**  
`docker run -e MY_VAR=MY_VAL ...`
  - **Files**  
`docker run -v /config.cf:/conf/config.cf ...` // Single file  
`docker run -v /config:/conf ...` // Directory of files
  - **CMD as arguments for ENTRYPOINT**  
`docker run image:latest '-Dswarm.project.stage=dev'`
- Docker provides CLI for managing secrets/configs (*Swarm only*)
- Kubernetes provides ConfigMaps/Secrets
  - Developers don't see configuration values and secrets anymore
  - Provides mechanisms for injecting configs and secrets into containers

# Application - Configure

- Applications can still package all configurations *(not recommended)*
  - controlled via switch (e.g.: via environment variable)
- Application should expose configurations
  - So that application can be configured for N stages without rebuild
- One build, one artifact, N stages
  - We want to move artifact or container over all stages
- Eclipse MicroProfile Config-API (<https://microprofile.io/>)
  - Abstracts developer from ConfigSource
  - Mechanisms to consume configurations as usually provided in the cloud
  - Several ConfigSource Types supported such as URL



# Application - Deploy

- Binary, Dockerfile and Scripts to build on target environment
- Ready to use Docker Image
  - Runs anywhere, where Docker is supported
- Docker Compose definition (*good for static application environments*)
  - docker-compose.yml, Binary and Dockerfile to build on target environment
  - or docker-compose.yml which references ready to use Docker Images
- Templates when running on CaaS/PaaS (*Kubernetes, Openshift, Azure, ...*)
  - Deploy provided Docker Image,
  - or deploy self built Docker Images (*BuildConfig - Openshift*)

# Application - Deploy

- When Dockerfile and Binaries are provided ensure
  - compatibility with used Linux Kernel,
  - compatibility with Docker Version,
  - compatibility of the scripts,
  - and the actuality and security of the provided resources.
- When Docker Images are provided ensure
  - to use a safe Docker Base-Image source (*RHEL*),
  - to keep Docker Images as small as possible (*RHEL Atomic, Alpine*),
  - to provide Docker Images via secured Docker Image-Registry,
  - that there are no secrets in the Docker Image layers,
  - and that Docker Container is ephemeral (*Drop and recreatable with little config*).



# Application - Backup

- Docker Container use Docker Volumes to keep data persistent
- It is not as easy as usual to get to the persistent data
- Prefer native backups over `docker container commit`
- Backup running container
  - `docker container exec -i -u root`  
`-v /backup:/backup`  
`mysql-db /usr/bin/mysqldump mydb > /backup/backup.sql`
- Backup stopped container
  - `docker container run -i -u root`  
`-v /backup:/backup --volumes-from mysql-db`  
`backup:latest /usr/bin/tar -zcvf /data/dump.tar.gz /mysql/data`

# Application - Restore

- If `docker container commit` was used, create Container of committed image
- Restoring running container
  - `docker container exec -i -u root`  
`-v /backup:/backup`  
`mysql-db /usr/bin/mysql < /backup/backup.sql`
- Restoring stopped container
  - `docker container run -i -u root`  
`-v /backup:/backup --volumes-from mysql-db`  
`backup:latest /usr/bin/tar -xvf /data/dump.tar.gz`
- Kubernetes/Openshift provide no native backup support

# Monitoring

- Docker CLI has little support for monitoring
  - `docker inspect <[Image-Id, Container-Id]>`
  - `docker logs -f <Container-Id>`
  - `docker stats <Container-Id>`
- Docker can contribute to Prometheus (*experimental*)
- Other tools available for monitoring
- Labels are crucial for monitoring containers (*Kubernetes, Openshift*)
- Application itself can contribute to any monitoring tool



# Monitoring - Logging

- stdin/stderr are captured by Docker
  - `docker logs -f <Container-Id>`
- Docker provides several drivers
- Application itself can send logs to log server
- No log to file, only stdout/stderr
- Applications log must provide transaction-id (*MicroProfile OpenTracing*)
- Logs are very important to analyze failures in a distributed system
- Openshift provides EFK stack (*Elastic, FluentD, Kibana*)

# Monitoring - Classical monitoring sys.

- Maybe they already support Docker/Kubernetes/Openshift
- Can, but doesn't have to run in a Docker Container
- New way of monitoring applications is via agent (*Java Agent*)
- Lot of provider already on the market:
  - CoScale Openshift Monitoring
  - Dynatrace Openshift Monitoring
  - hawt.io Java Web-Console (*Fuse Integration Services 2.0*)
  - ...

# Security - Docker Infrastructure

- Use build in Linux Security of Docker Host (*SELinux*)
- Restrict access Docker Host (*Linux user/group permissions*)
- Don't enable remote access (*Docker API-Server*)
  - If, then only via Client-Certificate-Authentication
- Don't run privileged containers (*no root access*)
- Don't use legacy repositories (*don't use registry v1*)
- Don't use insecure registries (*no Docker Hub*)



# Security - Docker Image-Registry

- Use encrypted transport protocol (*HTTPS only*)
- Restrict access to registry (*user/group permissions*)
- Restrict pushes to registry (*prevent distribution of insecure images*)
- Don't mirror insecure registries (*no Docker Hub*)
- Allow only signed content (*especially for production*)
- Separate registries for usage
  - Production (Released, signed and production ready Docker Images)
  - Tooling (Tooling for development)
  - ...

# Security - Docker Container

- Don't use unsigned Docker Images
- Don't use Docker Images from insecure/untrusted registries
- Don't use Docker Images which contain secrets/configurations
  - Provide secrets/configurations on startup
- Use minimized Docker Images (*no curl, ping or such installed*)
- Never run Docker Containers as root user
- Expose only necessary resources

# Security - Docker Container

- Security mostly applied outside, therefore:
  - Application uses <http://localhost:8080> instead of <https://localhost:443>
  - User access controlled via (OAuth2)
- Never run Docker Container as root user
- Expose only necessary resources
- Keep backend container in backend network (*no external access*)

